

COMP 790-124: Goals for today

- ▶ More notation: norms and inner products
- ▶ Constrained optimization and Lagrange multipliers
- ▶ Challenges in combining nonsmooth penalties
- ▶ Decomposing problems and Alternating Direction Method of Multipliers (ADMM)
- ▶ Fused Lasso illustration of ADMM

Norms

$$\ell_2 \quad \|\mathbf{x}\|_2 = \sqrt{\sum_i x_i^2}$$

$$\ell_1 \quad \|\mathbf{x}\|_1 = \sum_i |x_i|$$

$$\ell_\infty \quad \|\mathbf{x}\|_\infty = \max_i |x_i|$$

$$\text{Frobenius} \quad \|\mathbf{X}\|_F = \sqrt{\sum_i \sum_j x_{i,j}^2}$$

Using these we can write various costs more succinctly

$$\text{Linear Regression} \quad \frac{1}{2} \|\mathbf{y} - \beta_0 - \mathbf{X}\beta\|_2^2$$

$$\text{Ridge Regression} \quad \frac{1}{2} \|\mathbf{y} - \beta_0 - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_2^2$$

$$\text{Lasso} \quad \frac{1}{2} \|\mathbf{y} - \beta_0 - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1$$

$$\text{Elastic Net} \quad \frac{1}{2} \|\mathbf{y} - \beta_0 - \mathbf{X}\beta\|_2^2 + \lambda \|\beta\|_1 + \mu \|\beta\|_2^2$$

Bracket notation for inner products


Another convenient piece of notation

$$\mathbf{x}^T \mathbf{y} = \sum_i x_i y_i = \langle \mathbf{x}, \mathbf{y} \rangle$$

and then

$$\begin{aligned}\|\mathbf{x}\|_2^2 &= \langle \mathbf{x}, \mathbf{x} \rangle \\ \|\mathbf{x} - \mathbf{y}\|_2^2 &= \|\mathbf{x}\|_2^2 - 2 \langle \mathbf{x}, \mathbf{y} \rangle + \|\mathbf{y}\|_2^2\end{aligned}$$

Completing the square: couple of observations



$$\min_{\beta} f(\beta) + c = c + \min_{\beta} f(\beta)$$

Adding a constant to an objective changes the optimal value by c .

$$\operatorname{argmin}_{\beta} f(\beta) + c = \operatorname{argmin}_{\beta} f(\beta)$$

Adding a constant to an objective **does not** change the optimal β .

Completing the square

Suppose we have to solve a problem

$$\operatorname{argmin}_{\mathbf{x}} \underbrace{\frac{\rho}{2} \|\mathbf{x} - \mathbf{y}\|_2^2 + \langle \mathbf{x}, \mathbf{u} \rangle + f(\mathbf{x})}_A$$

then we can claim that the optimal \mathbf{x} is equal to

$$\operatorname{argmin}_{\mathbf{x}} \underbrace{\frac{\rho}{2} \left\| \mathbf{x} - \mathbf{y} + \frac{1}{\rho} \mathbf{u} \right\|_2^2}_B + f(\mathbf{x})$$

because difference between the objectives A and B is $-\frac{1}{2\rho} \|\mathbf{u}\|^2 + \langle \mathbf{u}, \mathbf{y} \rangle$ **a constant with respect to \mathbf{x} .**

Solving linear systems

Frequently problem of solving a linear systems of equations

$$\mathbf{Ax} = \mathbf{y}$$

is cast as optimization problem

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{Ax} - \mathbf{y}\|_2^2$$

And we already know one way of solving this problem.

Coordinate descent using following updates¹ :

$$x_i = \frac{\sum_j y_j^{-j} a_{i,j}}{\sum_j a_{i,j}}$$

This method is also known as Jacobi method.

¹ $y_i^{-j} = y_i - \sum_{k \neq j} a_{i,k} x_k$

Solving linear systems

Frequently problem of solving a linear systems of equations

$$\mathbf{Ax} = \mathbf{y}$$

is cast as optimization problem

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{Ax} - \mathbf{y}\|_2^2$$

and its solution can be obtained by setting gradient of objective to zero

$$\nabla \|\mathbf{Ax} - \mathbf{y}\|_2^2 = (\mathbf{Ax} - \mathbf{y})^T \mathbf{A} = \mathbf{0}$$

giving us **normal equations**

$$\mathbf{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}.$$

In MATLAB this is even more succinctly written as

$$\mathbf{x} = \mathbf{A} \backslash \mathbf{y}$$

Solving pairs of linear systems by stacking

Sometime we might run into an optimization problem

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{Ax} - \mathbf{y}\|_2^2 + \frac{\rho}{2} \|\mathbf{Bx} - \mathbf{z}\|_2^2$$

and this is equivalent to solving

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \left\| \begin{bmatrix} \mathbf{A} \\ \sqrt{\rho} \mathbf{B} \end{bmatrix} \mathbf{x} - \begin{bmatrix} \mathbf{y} \\ \sqrt{\rho} \mathbf{z} \end{bmatrix} \right\|_2^2$$

and in Matlab

$$\mathbf{x} = [\mathbf{A}; \text{sqrt}(\rho) * \mathbf{B}] \setminus [\mathbf{y}; \text{sqrt}(\rho) * \mathbf{z}]$$

Constrained optimization and Lagrange multipliers

The problems we looked at so far are unconstrained

$$\underset{\boldsymbol{\theta} \in \mathbf{R}^p}{\text{minimize}} \quad f(\boldsymbol{\theta}).$$

However sometimes we might want to impose constraints on our problems, for example

$$\begin{array}{ll} \underset{\boldsymbol{\theta} \in \mathbf{R}^p}{\text{minimize}} & f(\boldsymbol{\theta}) \\ \text{subject to} & g(\boldsymbol{\theta}) = 0. \end{array}$$

How do we solve such a problem?

Optimization of unconstrained objective

Given a differentiable objective $f(\boldsymbol{\theta})$ a critical point satisfies

$$\nabla f(\boldsymbol{\theta}) = \begin{bmatrix} \frac{\partial f(\boldsymbol{\theta})}{\partial \theta_1} \\ \frac{\partial f(\boldsymbol{\theta})}{\partial \theta_2} \\ \vdots \\ \frac{\partial f(\boldsymbol{\theta})}{\partial \theta_p} \end{bmatrix} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 0 \end{bmatrix} = \mathbf{0}_p.$$

One way to find $\boldsymbol{\theta}^*$ such that $\nabla f(\boldsymbol{\theta}^*) = \mathbf{0}_p$ is coordinate descent (ascent if we are maximizing instead of minimizing).

Optimization of a constrained problem

With constrained optimization problem we might not be able to find θ^* , such that $\nabla f(\theta^*) = \mathbf{0}_p$ and $g(\theta^*) = 0$.

Hence at optimum we might have to compromise by accepting $\nabla f(\theta^*) = \mathbf{v} \neq \mathbf{0}_p$ but in return $g(\theta^*) = 0$.

$$\nabla f(\theta^*) = \begin{bmatrix} \frac{\partial f(\theta^*)}{\partial \theta_1} \\ \frac{\partial f(\theta^*)}{\partial \theta_2} \\ \dots \\ \frac{\partial f(\theta^*)}{\partial \theta_p} \end{bmatrix} = \begin{bmatrix} v_1 \\ v_2 \\ \dots \\ v_p \end{bmatrix} = -\lambda \nabla g(\theta^*).$$

Intuitively, when tweaking θ_i^* , for each unit of improvement in f we pay λ in constraint violation.

All that remains is to set the price, λ , that is *sufficient* to make the constraint satisfied.

Optimization of constrained problem

A naive algorithm that solves a constrained problem

```
1:  $\lambda = 0$ 
2: repeat
3:    $\theta = \operatorname{argmin}_{\theta} f(\theta) + \lambda g(\theta)$ 
4:   if  $g(\theta) > 0$  then
5:      $\lambda = \lambda + \epsilon$ 
6:   end if
7:   if  $g(\theta) < 0$  then
8:      $\lambda = \lambda - \epsilon$ 
9:   end if
10: until  $|g(\theta)| < 10^{-12}$ 
```

The algorithm slowly adjusts the λ until the optimal solution satisfies the constraint.

Lagrangian

Given an optimization problem

$$\begin{array}{ll}\underset{\boldsymbol{\theta} \in \mathbf{R}^p}{\text{minimize}} & f(\boldsymbol{\theta}) \\ \text{subject to} & g_i(\boldsymbol{\theta}) = 0, i = 1, \dots, n \\ & h_j(\boldsymbol{\theta}) \leq 0, j = 1, \dots, m\end{array}$$

Lagrangian is a function

$$L(\boldsymbol{\theta}, \lambda, \mu) = f(\boldsymbol{\theta}) + \sum_{i=1}^n \lambda_i g_i(\boldsymbol{\theta}) + \sum_{j=1}^m \mu_j h_j(\boldsymbol{\theta})$$

with a requirement that $\mu_j \geq 0$.

The *prices* λ and μ are called **dual variables**, whereas $\boldsymbol{\theta}$ is a **primal variable**.

Lagrangian

More on Lagrangian's, primal and dual problem pairs, and more generally convex optimization in coming weeks.

For now, think of them as a means to get an unconstrained objective for a constrained problem.

An alternative for simple constraints

$$\begin{array}{ll}\text{minimize} & f(\theta) \\ \theta \in \mathbf{R}^p \\ \text{subject to} & \theta_i \geq 0, i = 1, \dots, n.\end{array}$$

We can rewrite by putting $\exp \gamma_i$ in place of θ_i

$$\begin{array}{ll}\text{minimize} & f(\exp \gamma) \\ \gamma \in \mathbf{R}^p \\ \text{subject to} & \exp \gamma_i \geq 0, i = 1, \dots, n.\end{array}$$

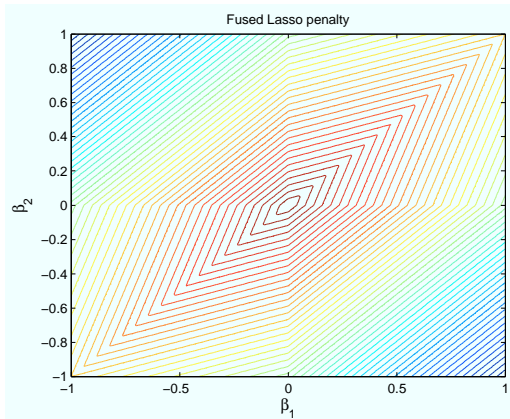
But, $\exp \gamma_i$ is always positive so the constraints are always satisfied and we can drop explicit constraints

$$\begin{array}{ll}\text{minimize} & f(\exp \gamma) \\ \gamma \in \mathbf{R}^p\end{array}$$

which is an unconstrained problem.

Coordinate ascent does not always work

An example of a problem where coordinate ascent gets stuck, fused lasso [2].



$$-\lvert\beta_1\rvert - \lvert\beta_2\rvert - 2\lvert\beta_1 - \beta_2\rvert$$

Optimization of challenging objectives - dual decomposition

Remove coupling across non-smooth objectives by adding more variables.

Challenging

$$\underset{\beta \in \mathbf{R}^2}{\text{minimize}} \quad |\beta_1| + |\beta_2| + \kappa |\beta_1 - \beta_2|$$

Easier after adding an auxiliary variable

$$\begin{array}{ll} \underset{\beta \in \mathbf{R}^2, \delta \in \mathbf{R}}{\text{minimize}} & |\beta_1| + |\beta_2| + \kappa |\delta| \\ \text{subject to} & \delta = \beta_1 - \beta_2 \end{array}$$

Making friends with Lagrangians

An example of a separable objective tied through a constraint

$$\begin{array}{ll}\underset{\mathbf{x}, \mathbf{y}}{\text{minimize}} & f(\mathbf{x}) + g(\mathbf{y}) \\ \text{subject to} & \mathbf{x} = \mathbf{y}\end{array}$$

and its Lagrangian

$$L(\mathbf{x}, \mathbf{y}, \lambda) = f(\mathbf{x}) + g(\mathbf{y}) + \lambda'(\mathbf{x} - \mathbf{y})$$

Augmented Lagrangian

$$AL(\mathbf{x}, \mathbf{y}, \lambda) = f(\mathbf{x}) + g(\mathbf{y}) + \sum_i \lambda_i(x_i - y_i) + \rho/2 \sum_i (x_i - y_i)^2$$

Alternating Direction Method of Multipliers[1] blueprint

$$AL(\mathbf{x}, \mathbf{y}, \lambda) = f(\mathbf{x}) + g(\mathbf{y}) + \sum_i \lambda_i (x_i - y_i) + \rho/2 \sum_i (x_i - y_i)^2$$

We will use k in superscript to denote state of variable after k^{th} iteration of algorithm.

$$\mathbf{x}^k = \underset{\mathbf{x}}{\operatorname{argmin}} AL(\mathbf{x}, \mathbf{y}^{k-1}, \lambda^{k-1})$$

$$\mathbf{y}^k = \underset{\mathbf{y}}{\operatorname{argmin}} AL(\mathbf{x}^k, \mathbf{y}, \lambda^{k-1})$$

$$\lambda_i^k = \lambda_i^{k-1} + \rho(x_i^k - y_i^k), i = 1, \dots, n$$

where $\rho > 0$.

ADMM for fused lasso on 2 variables

The fused lasso reformulated problem

$$\begin{aligned} & \underset{\beta \in \mathbf{R}^2, \delta \in \mathbf{R}}{\text{minimize}} && |\beta_1| + |\beta_2| + \kappa|\delta| \\ & \text{subject to} && \delta - (\beta_1 - \beta_2) = 0 \end{aligned}$$

and its Augmented Lagrangian

$$\begin{aligned} AL(\beta_1, \beta_2, \delta, \lambda) = & |\beta_1| + |\beta_2| + \kappa|\delta| + \\ & \lambda(\delta - (\beta_1 - \beta_2)) + \rho/2(\delta - (\beta_1 - \beta_2))^2 \end{aligned}$$

Alternating Direction Method of Multipliers for fused lasso on 2 variables

Iterate updates

$$\beta_1^k = \underset{\beta_1}{\operatorname{argmin}} AL(\beta_1, \beta_2^{k-1}, \delta^{k-1}, \lambda^{k-1})$$

$$\beta_2^k = \underset{\beta_2}{\operatorname{argmin}} AL(\beta_1^k, \beta_2, \delta^{k-1}, \lambda^{k-1})$$

$$\delta^k = \underset{\delta}{\operatorname{argmin}} AL(\beta_1^k, \beta_2^k, \delta, \lambda^{k-1})$$

$$\lambda^k = \lambda^{k-1} + \rho(\delta^k - (\beta_1^k - \beta_2^k)),$$

where $\rho > 0$.

ADMM for fused lasso on 2 variables

We will pause to note the update for β_1 :

$$\begin{aligned}\beta_1^k &= \operatorname{argmin}_{\beta_1} AL(\beta_1, \beta_2^{k-1}, \delta_1^{k-1}, \lambda^{k-1}) \\ &= \operatorname{argmin}_{\beta_1} |\beta_1| - \lambda^{k-1} \beta_1 + \rho/2 (\delta_1^{k-1} - (\beta_1 - \beta_2^{k-1}))^2\end{aligned}$$

can be cast into a form that we already know how to solve.

Completing squares

$$\operatorname{argmin}_{\beta_1} \underbrace{|\beta_1| - \lambda\beta_1 + \rho/2(\delta - (\beta_1 - \beta_2))^2}_{\text{A}}$$

is equal to

$$\operatorname{argmin}_{\beta_1} \underbrace{|\beta_1| + \rho/2(\beta_1 - (\delta + \beta_2 + \lambda/\rho))^2}_{\text{B}}$$

Why did we bother?

$$\operatorname{argmin}_{\beta_1} |\beta_1| + \rho/2(\beta_1 - (\delta + \beta_2 + \lambda/\rho))^2$$

Because this is a Lasso problem in a single variable.

Hence we obtain a closed-form update:

$$\beta_1^k = S(\delta^{k-1} + \beta_2^{k-1} + \lambda^{k-1}\rho, 1/\rho)$$


where

$$S(x, \lambda) \equiv \operatorname{sign}(x) \max(|x| - \lambda, 0)$$

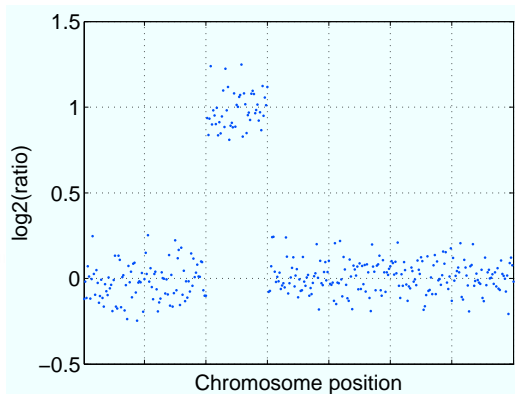
Key steps in optimizing coupled penalties

1. Reformulate problem to remove sharing of variables between non-smooth parts of objective
2. Write down Augmented Lagrangian for your reformulated problem
3. Iterate the ADMM scheme

Application of fused lasso to CNV data



U.S. National Library of Medicine



Fused Lasso Signal Approximator

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{y} - \mathbf{x}\|_2^2 + \lambda \|\mathbf{x}\|_1 + \mu \|\mathbf{D}\mathbf{x}\|_1$$

where \mathbf{y} is the vector we are trying to explain and \mathbf{D} is a matrix that ties different entries in \mathbf{x} together.

For example

$$\mathbf{D} = \begin{bmatrix} 1 & -1 & 0 & 0 & 0 \\ 0 & 1 & -1 & 0 & 0 \\ 0 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 1 & -1 \end{bmatrix}$$

gives term

$$\sum_{i=1}^3 |x_i - x_{i+1}|$$

Demos

Full derivation of ADMM for FLSA

Full derivation is available on the course webpage along with the code for demos.

We did ...

- ▶ More notation: norms and inner products
- ▶ Constrained optimization and Lagrange multipliers
- ▶ Challenges in combining nonsmooth penalties
- ▶ Decomposing problems and Alternating Direction Method of Multipliers (ADMM)
- ▶ Fused Lasso illustration of ADMM

Solving linear systems

Frequently problem of solving a linear systems of equations

$$\mathbf{Ax} = \mathbf{y}$$

is cast as optimization problem

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{Ax} - \mathbf{y}\|_2^2$$

and its solution can be obtained by setting gradient of objective to zero

$$\nabla \|\mathbf{Ax} - \mathbf{y}\|_2^2 = (\mathbf{Ax} - \mathbf{y})^T \mathbf{A} = \mathbf{0}$$

giving us **normal equations**

$$\mathbf{x} = (\mathbf{A}^T \mathbf{A})^{-1} \mathbf{A}^T \mathbf{y}.$$

In MATLAB this is even more succinctly written as

$$\mathbf{x} = \mathbf{A} \backslash \mathbf{y}$$

Solving pairs of linear systems by stacking

Sometime we might run into an optimization problem

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \|\mathbf{Ax} - \mathbf{y}\|_2^2 + \frac{\rho}{2} \|\mathbf{Bx} - \mathbf{z}\|_2^2$$

and this is equivalent to solving

$$\underset{\mathbf{x}}{\text{minimize}} \frac{1}{2} \left\| \begin{bmatrix} \mathbf{A} \\ \sqrt{\rho} \mathbf{B} \end{bmatrix} \mathbf{x} - \begin{bmatrix} \mathbf{y} \\ \sqrt{\rho} \mathbf{z} \end{bmatrix} \right\|_2^2$$

and in Matlab

$$\mathbf{x} = [\mathbf{A}; \text{sqrt}(\rho) * \mathbf{B}] \setminus [\mathbf{y}; \text{sqrt}(\rho) * \mathbf{z}]$$



Stephen Boyd, Neal Parikh, Eric Chu, Borja Peleato, and Jonathan Eckstein.

Distributed optimization and statistical learning via the alternating direction method of multipliers.

Foundations and Trends in Machine Learning, 3(1):1–122, 2011.



Robert Tibshirani, Michael Saunders, Saharon Rosset, Ji Zhu, and Keith Knight.

Sparsity and smoothness via the fused lasso.

Journal of the Royal Statistical Society Series B, 67(1):91–108, 2005.